

FE_scaling+balancing

November 8, 2021

In this notebook: * Feature engineering: scale the input variables between zero and 1 before applying logistic regression * balancing the dataset by under-sampling + LR * balancing the dataset by uover-sampling + LR * Balancing the dataset by a combination of under-sampling and over-sampling (SMOTE Tomek) + LR * Will use sklearn's LogisticRegression which is a regularized regression with hyperparameter "penalty" for regularization to address overfitting. Default is l2 (Ridge). * Embedded method of feature selection: Will try the "zeroing" regularization method l1 (LASSO).

0.0.1 Conclusions

- Balancing the dataset improved recall, but deteriorated accuracy and precision
- The method of resampling (undersampling, oversampling, SMOTETomek) performed equally

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, r
↪ recall_score, roc_curve, roc_auc_score
from sklearn import preprocessing

from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler
from imblearn.combine import SMOTETomek
```

/usr/local/lib/python3.7/dist-packages/sklearn/externals/six.py:31:

FutureWarning: The module is deprecated in version 0.21 and will be removed in version 0.23 since we've dropped support for Python 2.7. Please rely on the official version of six (<https://pypi.org/project/six/>).

"(<https://pypi.org/project/six/>).", FutureWarning)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:144:

FutureWarning: The sklearn.neighbors.base module is deprecated in version 0.22 and will be removed in version 0.24. The corresponding classes / functions should instead be imported from sklearn.neighbors. Anything that cannot be

```
imported from sklearn.neighbors is now part of the private API.  
warnings.warn(message, FutureWarning)
```

```
[2]: #Import Drive API and authenticate  
from google.colab import drive  
#Mount Drive to the Colab VM  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: #Load the dataset into pandas DataFrame  
df = pd.read_csv("/content/drive/MyDrive/Capstone_project/v2_credit_default.  
→csv")
```

```
[4]: #Seperate the independent and dependent variables.  
df_independent = df.drop(['Default'], axis=1)  
df_default = df['Default']
```

```
[5]: # split the data into 70% training and 30% test  
x_train, x_test, y_train, y_test = train_test_split(df_independent, df_default,   
→test_size=0.30, random_state=1)
```

Make sure the distribution of the dependent variable is the same in both training and test sets.

```
[6]: # training data  
y_train.value_counts().plot(kind='bar', title='Class distribution: Training_  
→data');  
y_train.value_counts()    #22% defaulters in the training data
```

```
[6]: 0    16354  
     1     4621  
     Name: Default, dtype: int64
```



```
[7]: # Test data
y_test.value_counts().plot(kind='bar', title='Class distribution: Test data')
y_test.value_counts() #22% defaulters in the test data
```

```
[7]: 0    6981
     1    2009
     Name: Default, dtype: int64
```



0.1 No resampling - used the original imbalanced dataset

```
[8]: # Scale input variables for training (x_train)
min_max_scaler = preprocessing.MinMaxScaler()
x_train_scaled = min_max_scaler.fit_transform(x_train)
```

0.1.1 Logistic regression

```
[9]: LR = LogisticRegression(random_state=1, solver='liblinear') #solver:
    ↳ 'lbfgs', 'newton-cg', 'liblinear', 'sag', 'saga', default penalty is 'l2'
# Fit the model
LR.fit(x_train_scaled, y_train)
# Predict using the scaled x_test
x_test_scaled = min_max_scaler.transform(x_test)
y_pred = LR.predict(x_test_scaled)
```


0.1.2 Performance metrics

```
[10]: cm = confusion_matrix(y_test, y_pred)
print(cm)
print('accuracy', accuracy_score(y_test, y_pred))
print('precision', precision_score(y_test, y_pred))
print('recall', recall_score(y_test, y_pred))
```


```
[[6784  197]
 [1527  482]]
accuracy 0.8082313681868744
precision 0.7098674521354934
recall 0.23992035838725734
```

```
[11]: # Graphic visualization of the confusion matrix
ax= plt.subplot()
sns.heatmap(cm,annot=True,fmt='g', ax=ax, cmap='Greens'); #annot=True to
    →include the numeric values, ftn='g' to avoid scientific notation

ax.set_xlabel('Predicted');ax.set_ylabel('Actual');
ax.xaxis.set_ticklabels(['NoDefault', 'Default']);
ax.yaxis.set_ticklabels(['NoDefault', 'Default']);
```



```
[12]: # ROC curve, and ROC AUC
y_pred_proba = LR.predict_proba(x_test_scaled)[::,1] #probability estimated
↳from LR #[start:stop:step]
FPR, TPR, _ = roc_curve(y_test, y_pred_proba) #roc_curve(y_true,
↳y_score)
auc = roc_auc_score(y_test, y_pred_proba)
plt.plot(FPR,TPR,label="LR-no resampling, auc="+str(auc))
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC curve')
plt.legend(loc=4) #place the legend in the lower right corner
plt.show()
```



0.2 Under-sampling the majority class + train a LR model

```
[13]: rus = RandomUnderSampler(return_indices=True, random_state=1)
X_rus, y_rus, id_rus = rus.fit_resample(x_train_scaled, y_train.squeeze())
#Now to check the training data distribution after undersampling
pd.Series(y_rus.reshape(-1)).value_counts().plot(kind='bar', title='Training
↳data - undersampling'); #convert numpy array to panda's series to use
↳panda's value_counts()
#training data is now balanced
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
```




```
[14]: #repeat fitting and predicting with logistic regression:
LR_us = LogisticRegression(random_state=1, solver='liblinear')
# Fit the model
LR_us.fit(X_rus, y_rus.squeeze())
# Predict using the scaled x_test
y_pred = LR_us.predict(x_test_scaled)

[15]: cm = confusion_matrix(y_test, y_pred)
print(cm)
print('accuracy', accuracy_score(y_test, y_pred))
print('precision', precision_score(y_test, y_pred))
print('recall', recall_score(y_test, y_pred))
# Graphic visualization of the confusion matrix
ax= plt.subplot()
sns.heatmap(cm,annot=True,fmt='g', ax=ax, cmap='Greens'); #annot=True to
→include the numeric values, ftn='g' to avoid scientific notation

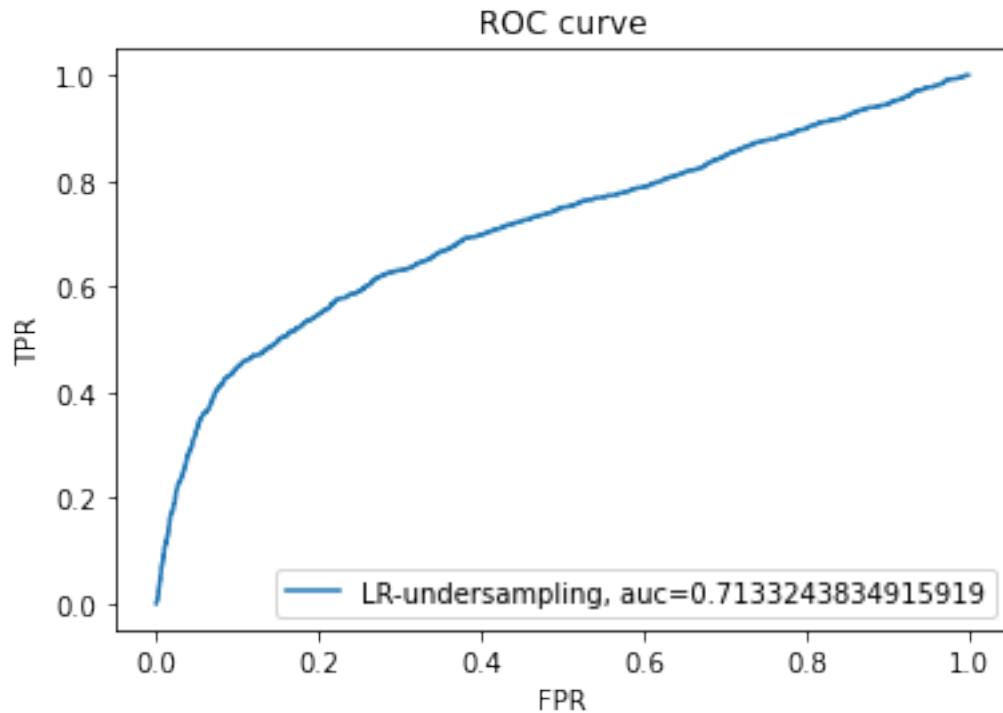
ax.set_xlabel('Predicted');ax.set_ylabel('Actual');
ax.xaxis.set_ticklabels(['NoDefault', 'Default']);
ax.yaxis.set_ticklabels(['NoDefault', 'Default']);
```

```
# result: with undersampling the training data, accuracy and precision  
→ suffered, but recall improved
```

```
[[5047 1934]  
 [ 764 1245]]  
accuracy 0.699888765294772  
precision 0.39163258886442276  
recall 0.6197112991538078
```



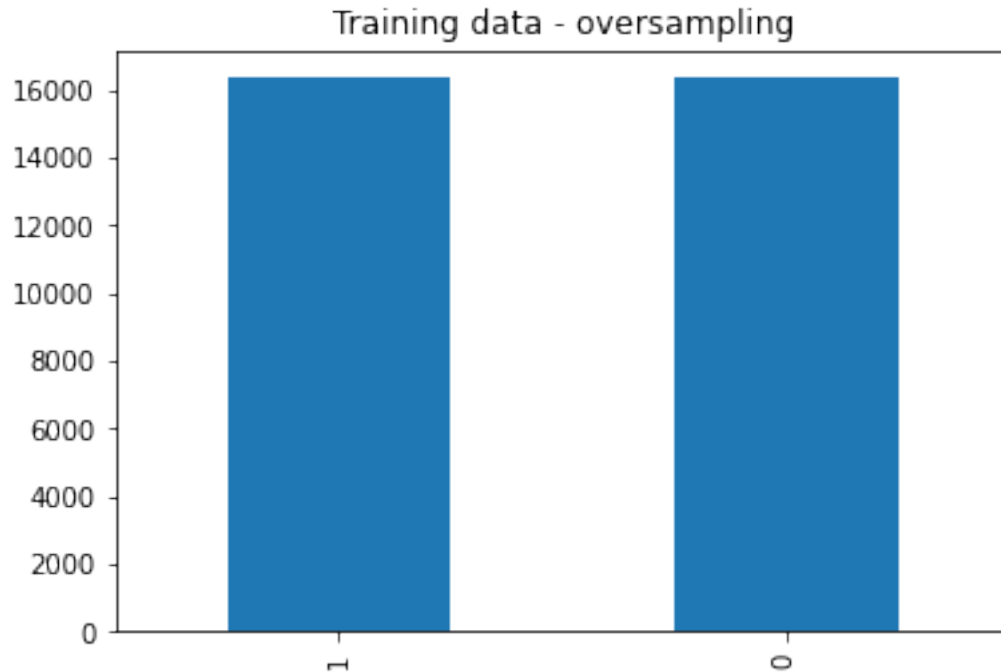
```
[16]: # ROC curve, and ROC AUC  
y_pred_proba = LR_us.predict_proba(x_test_scaled)[:,:1] #probability estimated  
→ from LR #[start:stop:step]  
FPR, TPR, _ = roc_curve(y_test, y_pred_proba) #roc_curve(y_true,  
→ y_score)  
auc = roc_auc_score(y_test, y_pred_proba)  
plt.plot(FPR,TPR,label="LR-undersampling, auc="+str(auc))  
plt.xlabel('FPR')  
plt.ylabel('TPR')  
plt.title('ROC curve')  
plt.legend(loc=4)  
plt.show()
```

0.3 Over-sampling the majority class + train a LR model

```
[17]: ros = RandomOverSampler(return_indices=True, random_state=1)
X_ros, y_ros, id_ros = ros.fit_resample(x_train_scaled, y_train.squeeze())
#Now to check the training data distribution after oversampling
pd.Series(y_ros.reshape(-1)).value_counts().plot(kind='bar', title='Training_
↳data - oversampling'); #convert numpy array to panda's series to use_
↳panda's value_counts()
#training data is now balanced
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
```



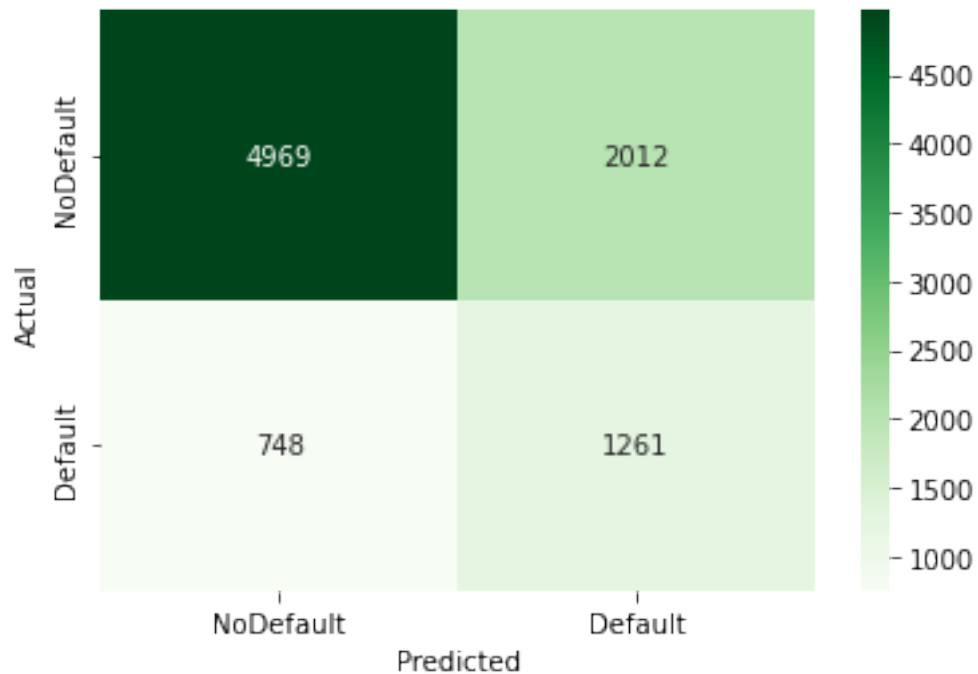
```
[18]: #repeat fitting and predicting with logistic regression:
LR_os = LogisticRegression(random_state=1, solver='liblinear')
# Fit the model
LR_os.fit(X_ros, y_ros)
# Predict using the scaled x_test
y_pred = LR_os.predict(x_test_scaled)

[19]: cm = confusion_matrix(y_test, y_pred)
print(cm)
print('accuracy', accuracy_score(y_test, y_pred))
print('precision', precision_score(y_test, y_pred))
print('recall', recall_score(y_test, y_pred))
# Graphic visualization of the confusion matrix
ax= plt.subplot()
sns.heatmap(cm,annot=True,fmt='g', ax=ax, cmap='Greens'); #annot=True to
    ↳include the numeric values, ftm='g' to avoid scientific notation


ax.set_xlabel('Predicted');ax.set_ylabel('Actual');
ax.xaxis.set_ticklabels(['NoDefault', 'Default']);
ax.yaxis.set_ticklabels(['NoDefault', 'Default']);
# result: oversampling has similar performance to undersampling
```

```
[[4969 2012]
 [ 748 1261]]
accuracy 0.692992213570634
```

precision 0.3852734494347693
recall 0.6276754604280737




```
[20]: # ROC curve, and ROC AUC
y_pred_proba = LR_os.predict_proba(x_test_scaled)[:,:1]
FPR, TPR, _ = roc_curve(y_test, y_pred_proba) #roc_curve(y_true, y_score)
auc = roc_auc_score(y_test, y_pred_proba)
plt.plot(FPR,TPR,label="LR-over sampling, auc="+str(auc))
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC curve')
plt.legend(loc=4)
plt.show()
```



0.4 Using both undersampling and oversampling + train a LR model

```
[21]: x_smt, y_smt = SMOTETomek(random_state=1).fit_sample(x_train_scaled, y_train.
    ↪squeeze())
    #Now to check the training data distribution:
    pd.Series(y_smt.reshape(-1)).value_counts().plot(kind='bar', title='Training_
    ↪data - SMOTETomek');
    #training data is now balanced
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function safe_indexing is deprecated; safe_indexing is deprecated
in version 0.22 and will be removed in version 0.24.
    warnings.warn(msg, category=FutureWarning)
```



```
[22]: #repeat fitting and predicting with logistic regression:
LR_smt = LogisticRegression(random_state=1, solver='liblinear', penalty='l1')
    ↳ #default penalty='l2'
# Fit the model
LR_smt.fit(x_smt, y_smt)
# Predict using the scaled x_test
y_pred = LR_smt.predict(x_test_scaled)
```


```
[34]: #same performance between l1 and l2 penalty. But the l1 penalty reduced some of
    ↳ the coefficients to zero.
column_labels = df_independent.columns.tolist()
coef = LR_smt.coef_.squeeze().tolist()
# Zip together
labels_coef = list(zip(column_labels, coef))
print(labels_coef)
print(LR_smt.intercept_)
# coefficients for Repay_June, and Bill_May are reduced to zero.
```

```
[('LIMIT_BAL', -0.9911510708551642), ('SEX', -0.11380413540340302),
('EDUCATION', -0.28234179240855295), ('MARRIAGE', -0.3848891410296028), ('AGE',
0.31170516619952787), ('Repay_Sept', 5.367874181334451), ('Repay_Aug',
1.0787473816376605), ('Repay_July', 0.6136194082941878), ('Repay_June', 0.0),
('Repay_May', 0.04802380184376096), ('Repay_Apr', 0.04667828550850049),
('Bill_Sept', -4.6052808076736635), ('Bill_Aug', 0.8594555487875079),
('Bill_July', 1.1474550781100228), ('Bill_June', 1.3174096973112688),
```

```
('Bill_May', 0.0), ('Bill_Apr', 0.8364492020336722), ('Pay_Sept',  
-15.373818503130556), ('Pay_Aug', -10.944389684851998), ('Pay_July',  
-6.691640093699221), ('Pay_June', -3.208914878204537), ('Pay_May',  
-1.308362184597507), ('Pay_Apr', -1.6037870340479032)]
```

```
[24]: cm = confusion_matrix(y_test, y_pred)  
print(cm)  
print('accuracy', accuracy_score(y_test, y_pred))  
print('precision', precision_score(y_test, y_pred))  
print('recall', recall_score(y_test, y_pred))  
# Graphic visualization of the confusion matrix  
ax= plt.subplot()  
sns.heatmap(cm,annot=True,fmt='g', ax=ax, cmap='Greens');  
  
ax.set_xlabel('Predicted');ax.set_ylabel('Actual');  
ax.xaxis.set_ticklabels(['NoDefault', 'Default']);  
ax.yaxis.set_ticklabels(['NoDefault', 'Default']);  
# result: similar performance to under and over sampling
```

```
[[4822 2159]  
 [ 719 1290]]  
accuracy 0.6798665183537264  
precision 0.3740214554943462  
recall 0.6421105027376804
```



```
[25]: # ROC curve, and ROC AUC
y_pred_proba = LR_smt.predict_proba(x_test_scaled)[:,:1]
FPR, TPR, _ = roc_curve(y_test, y_pred_proba) #roc_curve(y_true, y_score)
auc = roc_auc_score(y_test, y_pred_proba)
plt.plot(FPR,TPR,label="LR-SMOTETomek, auc="+str(auc))
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC curve')
plt.legend(loc=4)
plt.show()
```

